

Testing of Serializability

A serializability graph is used to test the serializability of a schedule.

For testing of serializability it is simple and efficient method is to construct a directed graph called a precedence graph from

$$G = (V, E)$$

where V consists of set of vertices and E consists of a set of edges. The set of vertices is used to contain all the transactions participating in the schedule.

The set of edges is used to contain all edges $T_i \rightarrow T_j$ for which one of the three conditions holds

1. Create a node $T_i \rightarrow T_j$ if T_i executes write (Q) before T_j executes read (Q).
2. Create a node $T_i \rightarrow T_j$ if T_i executes read (Q) before T_j executes write (Q).
3. Create a node $T_i \rightarrow T_j$ if T_i executes write (Q) before T_j executes write (Q).

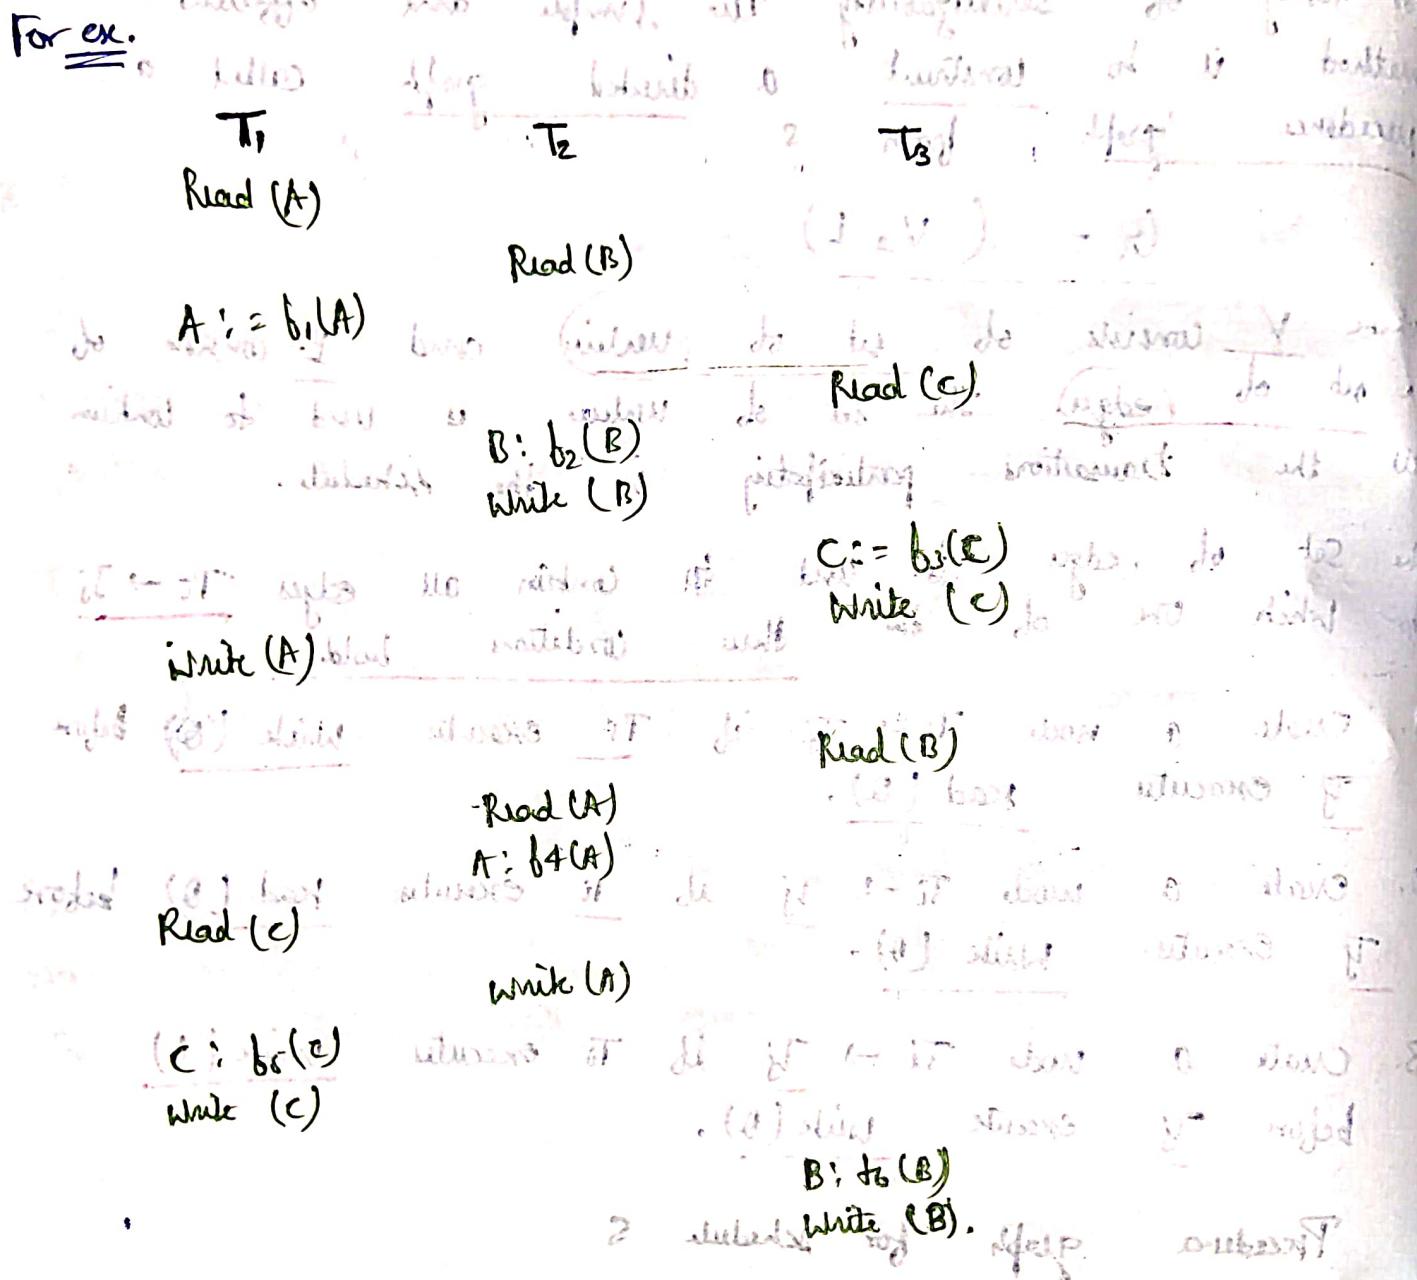
Precedence graph for schedule S



- b) A precedence graph contains a single edge $T_i \rightarrow T_j$ if all the instructions of T_i are executed before the first instruction of T_j is executed.

↳ A precedence graph for schedule S_1 contains a cycle. If the schedule is non-serializable, then the precedence graph has no cycle, then it is known as serializable.

For ex.



Schedule S_1

Explanation: T_1 gets applied a initial flag marking A

$\text{Read}(A)$ In T_1 , no subsequent writes to A, so no new edges.

$\text{Read}(B)$

In T_2 , no subsequent writes to B, so no new edges.

Read (C)

↳ In T_3 , no subsequent write to C, so no new edges

Write (D)

↳ D is subsequently read by T_3 , so add edge $T_2 \rightarrow T_3$

Write (C)

↳ C is subsequently read by T_1 , so add edge $T_2 \rightarrow T_1$

Write (A)

↳ A is subsequently read by T_1 , so add edge $T_1 \rightarrow T_2$

Write (A)

↳ In T_2 no subsequent reads to A, so no new edges

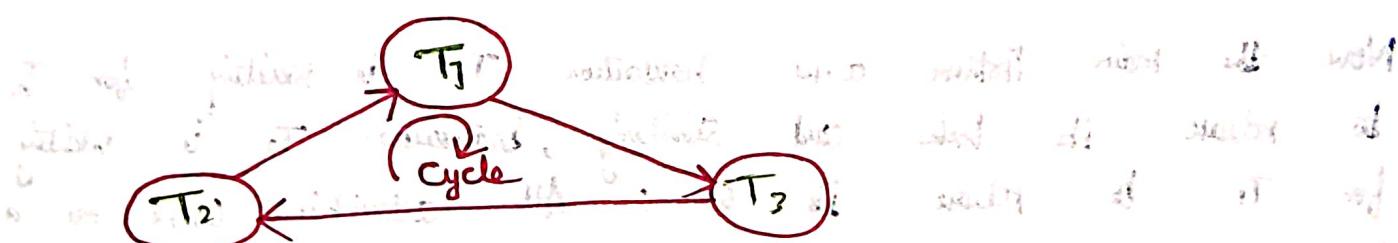
Write (C)

↳ In T_1 , no subsequent read of Block A, no new edges

Write (B)

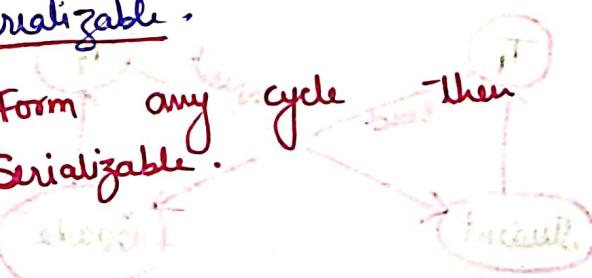
↳ In T_2 no subsequent read to Block B, no new edges

Precedence graph for schedule S_1 :



The precedence graph for schedule S_1 contains a cycle. That's why schedule S_1 is non serializable.

When it does not form any cycle then the schedule is Serializable.

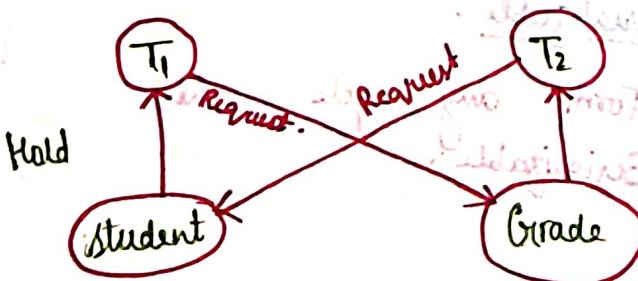


Deadlock in DBMS.

- ↳ A deadlock is a condition where two or more transaction are waiting indefinitely for one another to give up lock.
- ↳ A system is said to be in deadlock state when a set of transaction, Every transaction is waiting for another transaction to finish its operation.
- ↳ A deadlock is a state of statements that may result when two or more transaction are each waiting for locks held by the other to be released.

For Example: In the student table transaction T₁ holds a lock on some rows and needs to update some rows in the grade table. Simultaneously, transaction T₂ holds locks on some rows in the grade table and needs to update the rows in the student table held by transaction T₁.

Now the main problem arises Transaction T₁ is waiting for T₂ to release its locks and similarly, transaction T₂ is waiting for T₁ to release its lock. All activities come to a halt state and remain at a standstill.



①

"Jai Shri Ram"

Aspect	1NF	2NF	3NF	BCNF	4 NF	5 NF
Focus	Eliminates repeating groups	Removes transitive dependency	Ensures every determinant is a candidate key.	Ensures every determinant is a candidate key.	Ensures every dependency is implied by candidate keys.	Ensures every dependency is implied by candidate keys.
Attribute Requirement	Atomic Attribute	No Partial dependency on any candidate key.	No Transitive dependency on any candidate key.	No trivial functional dependency is o candidate key.	No multi valued dependency is o candidate key.	Each 'join' dependency is implied by candidate keys.
Partial Dependency	Allowed	Not Allowed	Not Allowed	Not Allowed	Not Allowed	Not Allowed.
Transitive dependency	Allowed	Allowed	Allowed	Not Allowed	Not Allowed	Not Allowed.

Multi-valued dependency	join dependency	Allowed	Allowed	Not Allowed	Not Allowed.
Example scenario	Table with non-key column group, is dependent on primary key column	Table with non-trivial FD's	Table with no trivial FD's	Table with non-trivial join dependency.	Table with no join dependency.
Notes	Additional simplest form focus on atomicity.	Build on 1NF of eliminating partial dependency.	Build on 2NF, stronger version of 3NF.	Ensures tables are decomposed without loss of data.	Ensures tables are decomposed without loss of data.

Q. Is relational model what do you mean by Cardinality?

Ans. In the relational model of DB. Cardinality refers to the no. of elements or the no. of rows in a set of different contents.

↳ Cardinality refers to two different contents.

I. Cardinality of a Table
↳ This refers to the no. of rows (tuple) in a table.
For eg. If a table Students has 100 rows then the cardinality of the table is 100.

II. Cardinality of Relationships.
↳ This refers to the uniqueness of data values contained in a particular column (attribute) of a table and more commonly to the relationships between tables.

↳ The relationship cardinality specifies the no. of instances of one entity that can or must be associated with each instance of another entity.

They types of Cardinality

- (i) One to One (1:1)
- (ii) One to Many (1:N)
- (iii) Many to One (N:1)
- (iv) Many to Many (N:N)

Q3 How can you map a conceptual model to a relational model?

Ans. Mapping a conceptual model (often represented as an Entity - Relationship Diagram, or ERD) to a relational model involves converting entities, relationships and attributes into tables columns and keys.

Step by step mapping

1. entities to tables.

↳ Each entity in the conceptual model becomes a table in the relational model.

↳ attributes of the entity become columns.

↳ PK of the entity becomes PK of the table.

2. Attributes to Column

↳ Attribute to column.

↳ includes appropriate data types for each column.

↳ If an attribute is a multi valued attribute

create a separate table for it.

3. Relationship to Foreign Keys

→ One to one Relationship

↳ Add the PK of one entity as foreign key in the related entity table.

→ Many to one relationship

↳ One to many relationship

↳ Add the PK of the "one" side entity as a foreign key in the "many" side entity's table.

4. Many to Many

↳ Create a new table to represent the relationship.

↳ This new table will have foreign keys referencing the PKs of the selected entities.

↳ PK of this new table is considered composite key made up of two foreign keys.

4. Weak entities

↳ Convert weak entities into tables.

↳ Create a FK column that references the PK of the primary entity table.

↳ Include attributes of the tables and its

The pk of weak entity table be typically a composite key consisting of its own identifier and the foreign key.

5. Composite Attribute

- ↳ Break down of composite attributes into individual attributes and map them to separate columns.

6. Multi - Valued Attributes

- ↳ Create a new table for multi valued attribute.
 - ↳ This new table will include a FK column referencing the pk of the original entity.
 - ↳ The pk of the new table will usually be a composite key made up of the FK and the multi valued attribute.

~~Q4~~ list two reasons why we may choose to define a view.

Ans.

1. Data security and Access control

↳ Views can be used to restrict access to specific data in a DB, thereby enhancing security. By granting rather than Control what data manipulate.

For instance

↳ Sensitive data protection

↳ You can create a view that excludes sensitive columns (e.g. salary information) so users can access non-sensitive data without compromising privacy.

↳ Simplified access control

↳ Instead of setting complex permissions on each table, grant access to views which encapsulate multiple permission in one place.

E. Simplification and Abstraction of complex Queries.

↳ Views can simplify complex queries by encapsulating them, making it easier for user to retrieve and manipulate data without having to understand or work with complex SQL.

This abstraction can lead to more readable and maintainable code.

Q5. What is Phantom Phenomenon?

Ans. It is a specific type of concurrency anomaly that can occur in DB systems which transactions are executed concurrently. It happens when a transaction reads a set of rows that satisfy a certain condition, but a subsequent read within the same transaction returns additional rows that were inserted by another concurrent transaction.

Q8. Which protocol always ensures recoverable schedule?

Aw. To ensure a recoverable schedule in db system protocols are designed to manage transaction concurrency and maintain stated consistency. There should be such protocol. at least one suitable will be such protocol.

List. Strict Two-Phase locking (strict 2PL) Protocol -

↳ strict 2PL is a specific type of the two Phase locking (2PL) protocol that ensures a recoverable schedule by imposing additional constraints on how locks are released.

How it works:-

(i) Lock Acquisition Phase :-

During this phase a transaction can acquire any no. of locks (read or write) it needs, but if it cannot release any locks.

(ii) Lock Release Phase

Once a transaction releases its first lock, it cannot acquire any new locks.

↳ In the strict version of 2PL, all locks (both read and write) held by a transaction are only released when the transaction is committed or aborted. This ensures that no other transaction can read or write any data modified by the transaction until it is finished, preventing cascading rollbacks and ensuring a recoverable schedule.

- * Properties of strict 2PL.
- ↳ Prevents Cascading Rollbacks
 - ↳ Because locks are only released at the end of a transaction, no other transaction can access the modified data until the transaction is complete. This ensures that if a transaction fails, no other transaction need read between the uncommitted data, avoiding cascading rollbacks.
 - ↳ Ensure Recoverable Schedule
 - ↳ Since transactions cannot read uncommitted data, any schedule produced under strict 2PL will be recoverable. A transaction can only commit if all the changes it has read have also committed.

Shared lock → read
 Exclusive lock → write
 pre changing lock

Snapshot lock

Two phase lock
 Start Two phase lock..

Q7 What is the possible violation if a application program use isolation level repeatable read

Ans: using the Repeatable Read - isolation level in a DB system provides a higher level of consistency compared to lower isolation levels like Read committed but however it is not with the highest level of isolation and certain type of anomalies can still occur. Specifically the main possible violation that can occur under Repeatable Read is the Phantom Read anomaly.

and here we can see that two rows are inserted and the newly inserted rows are not visible to the user.

Programmer can use the snapshot isolation level to avoid this problem.

Snapshot isolation level is a hybrid of both row and transaction level.

Snapshot isolation level is used to avoid phantom reads.

Snapshot isolation level is used to avoid dirty reads.

Snapshot isolation level is used to avoid non-repeatable reads.

Q.8. Explain with the help of examples the concept of insertion anomalies and deletion anomalies.

Ans. Insertion anomalies and deletion anomalies are type of data integrity issue that can occur in relational DB, particularly when the DB schema is not properly normalized.

The anomalies can lead to redundant, inconsistent or incomplete data. Here are

↳ Insertion Anomalies

↳ It occurs when certain attribute cannot be inserted into the DB without the presence of other attributes. This often happens when data is stored in a single table that combines multiple entities.

Ex:	Student ID	Student Name	Course ID	Course Name	Instructor Name
	1	Alice	C101	Meth	Prof. Smith
	2	Bob	C102	Science	Prof. Jones.

If we want to insert a new data with course ID given course name "Math" → Then we don't have Instructor No " " to put the data in the table.

We might insert a dummy student, leading to incorrect or misleading data.

↳ Deletion Anomalies

↳ It occurs when the deletion of data representing one fact necessitates the deletion of data representing another fact, which should not be deleted. This often results in loss of valuable information.

e.g. StudId.

	Std. Name	CourseId	CourseName	InstructorName
1	Alice	C101	Math	Prof. Smith
2.	Bob	C102	Science	Prof. Jones
3	Charlie	C101	Math	Prof. Smith

If we want to delete the row with course Id C101 then it also removes both the rows. But we want to remove only one.

This would result in loss of information.

Q9 Solution to Anomalies

Normalization is the process of organizing data in a DB to reduce redundancy and improve data integrity. It involves decomposing tables into smaller, related tables and defining relationships b/w them.

We can normalize the Student Course table into two separate tables.

1. Student Table :		Student ID	Student Name
1.	John	1	Alice
2.	David	2	Bob
3.	Will	3	Charlie
4.	Sam	4	David

2. Course Table :		Course ID	Course Name	Instructor Name
1.	Math	C101	Math 101	Prof. Sciume
2.	Science	C102	Science	Prof. Jones
3.	History	C103	History	Prof. Brown
4.	English	C104	English	Prof. Green

3. Enrollment Table :		Student ID	Course ID
1.	John	1	C101
2.	Bob	2	C102
3.	Charlie	3	C101
4.	David	4	C103

Benefits

- Insertion Anomaly Prevention:
↳ we can now add a new course without needing to enroll a student into it.

inserting a row into the enrollment table does not affect the course table.

- Deletion Anomaly Prevention:
we can delete a student's enrollment entry without losing any information about the course. For instance, deleting Charlie's enrollment (student ID = 3, course ID = C101) from the enrollment table does not affect the course table.

Q10. What is the goal of query optimization? Why is optimization important?

Ans. :- The goal of query optimization is to improve the performance and efficiency of DB queries, ensuring that they execute quickly and with few resources as possible while still producing the correct results.

Importance of query optimization

1. Improved Performance

↳ Efficient queries result in faster response time, reducing latency and improving overall system performance.

This is especially critical in applications where users expect near-instantaneous responses, such as web apps or real-time analytics systems.

2. Resource Utilization

↳ Optimized queries consume fewer system resources such as CPU, memory, and disk I/O.

This leads to better resource utilization, allowing the system to handle more concurrent requests and users without degradation in performance.

3. Scalability

As the volume of data and user requests increase, a well optimized DB can scale more effectively. By reducing the

the computational burden of each query, if optimization helps the system maintain performance level even under of heavy load.

* Cost Reduction → In cloud computing environments where resources are billed based on usage, query optimization can lead to cost savings by reducing the amount of resources consumed. This is particularly important for organizations managing large-scale data operations.

* User Experience → Faster query response time improves the user experience by providing more responsive applications and reducing wait times. This can lead to higher user satisfaction and engagement.

* Complexity Management → Query optimization often involves analyzing and understanding the underlying data model, indexes, and query execution plan. This process can uncover inefficiencies or bottlenecks in the db design, leading to improvement in overall system architecture and complexity management.

* Adaptability → Optimization techniques evolve with advancements in db technology and hardware capabilities.